

Chapter 3

Part 1 DATABASE RECOVERY TECHNIQUES



KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN



Outline

1. Goals
2. Type of Failures
3. Major Concepts
 - Transaction log
 - Checkpoint
4. Database Recovery Techniques





Example

- ☐ Accounts: A : 1000 \$, B : 2000 \$
- ☐ Transfer 50 \$ from A to B.
- ☐ Let's consider the state:
 - ☐ A had updated: $A := A - 50$
 - ☐ B had not been updated: $B := B + 50$
 - ☐ And there was power cut!
- ☐ When the system restarted:
 - ☐ If T was executed again: $A = 900$
 - ☐ If not: $A = 950$ and $B = 2000$.
 - ☐ How should the system be recovered?





Goals

- Database Recovery is the process of restoring the database and the data to a consistent state. This may include restoring lost data up to the point of the event (e.g. system crash).
- This is done by RM – the Recovery Manager.
- Automatic database recovery helps saving manual effort after every crash.





Goals

- Transaction is the basic unit of Database Recovery.
- Among ACID properties of Transaction, Database Recovery helps retain Atomic and Durability.





Types of failures

- Transaction failures
 - Rolling back caused by Deadlock or as requested by the scheduler.
 - During failure, system still operates normally.
 - Frequency: some times/minute.
- System failures
 - System can no longer operate. The reason may be failures in the processing unit, power cut or software failure.
 - Only cause data lost on RAM.
 - Frequency: some times/month.





Types of failures

- Media failure
 - Ex: HDD crash.
 - Cause the lost of a part or entire database.
 - Frequency: some times/year.
- Software failures
 - Logical error of applications which access database, leading to transaction failure.



Note

- Regardless of failure reasons, we have to consider 2 issues:
 1. Data lost on database buffer.
 2. Data lost on storage media.





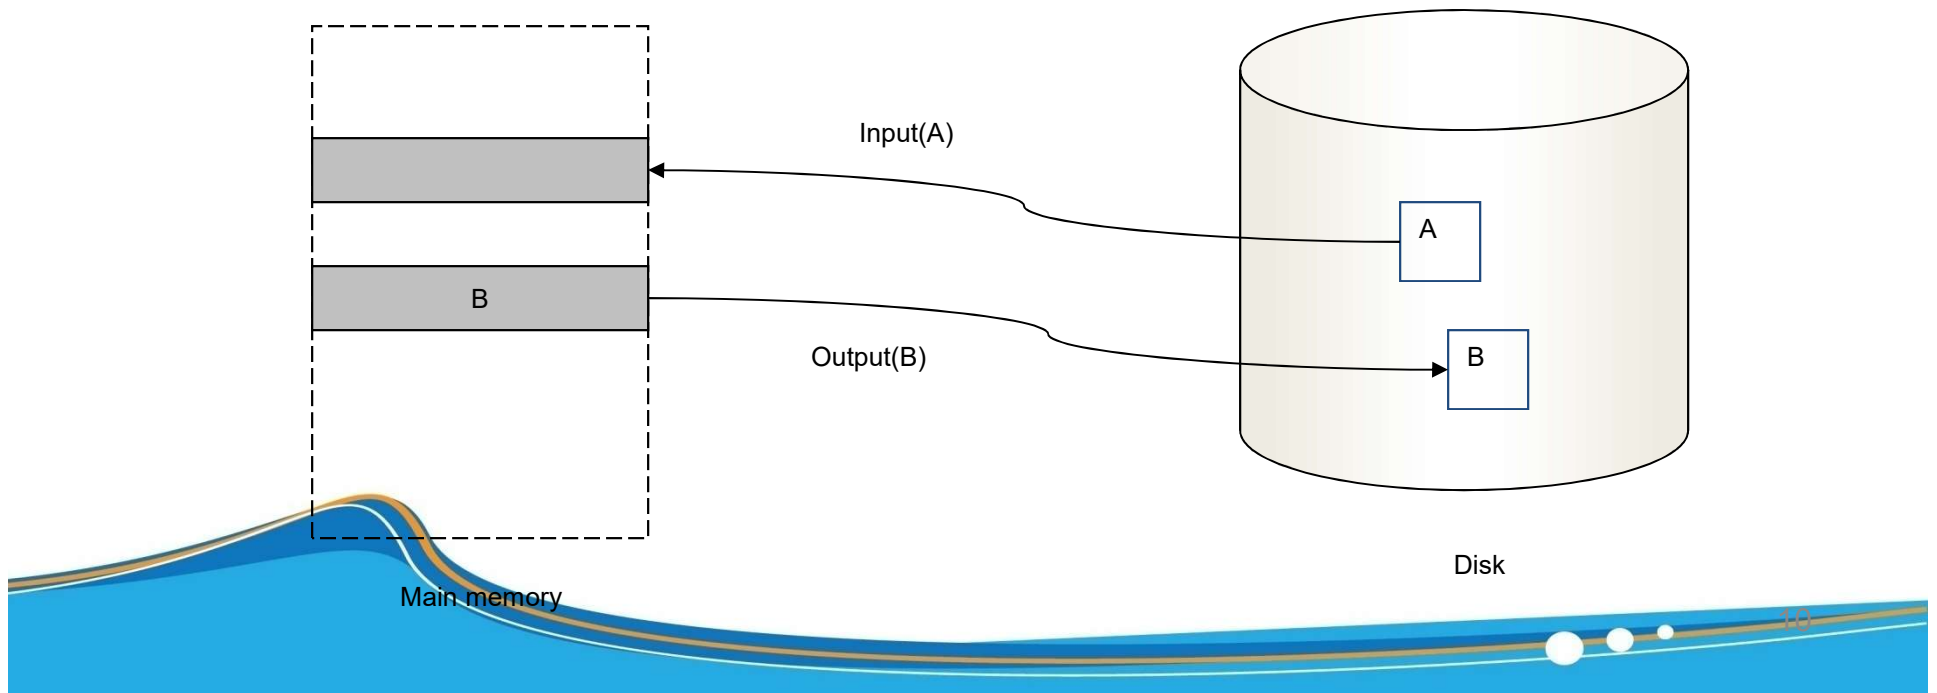
Backup

- DBMS provide Backup mechanism.
 - Full backup / Differential backup.
 - Handle data lost on storage media.



DB access

- Data is read or written in primary unit called “block”.
- ▣ Physical block: Stored in Media.
- ▣ Buffer block: Stored in buffer.



DB access

- Read (X): assign X to local variable
 - ▣ If the data block with X is not in buffer, then Input (X).
 - ▣ assign X to local variable Xi.
- Write (X): assign Xi to X
 - ▣ If the data block with X is not in buffer, then Input (X)..
 - ▣ Gán giá trị xi cho X (trên buffer block có chứa X).





Buffer management

- Buffer:
 - Dữ liệu mất khi có sự cố hệ thống.
 - Không gian hạn chế.
- Chiến lược thay thế để định ra vùng trống trên buffer dùng để nạp dữ liệu mới.
 - FIFO.
 - LRU.





Note

□ Database has 2 parts:


- Physical Database and
- Database Buffer

□ Log has 2 parts:

- Physical Log and
- Log Buffer

□ Failure \Rightarrow data lost in database buffer \Rightarrow use log file to recover data.

□ Failure \Rightarrow data lost in log buffer \Rightarrow redo DB manipulation which has not been saved to physical DB.



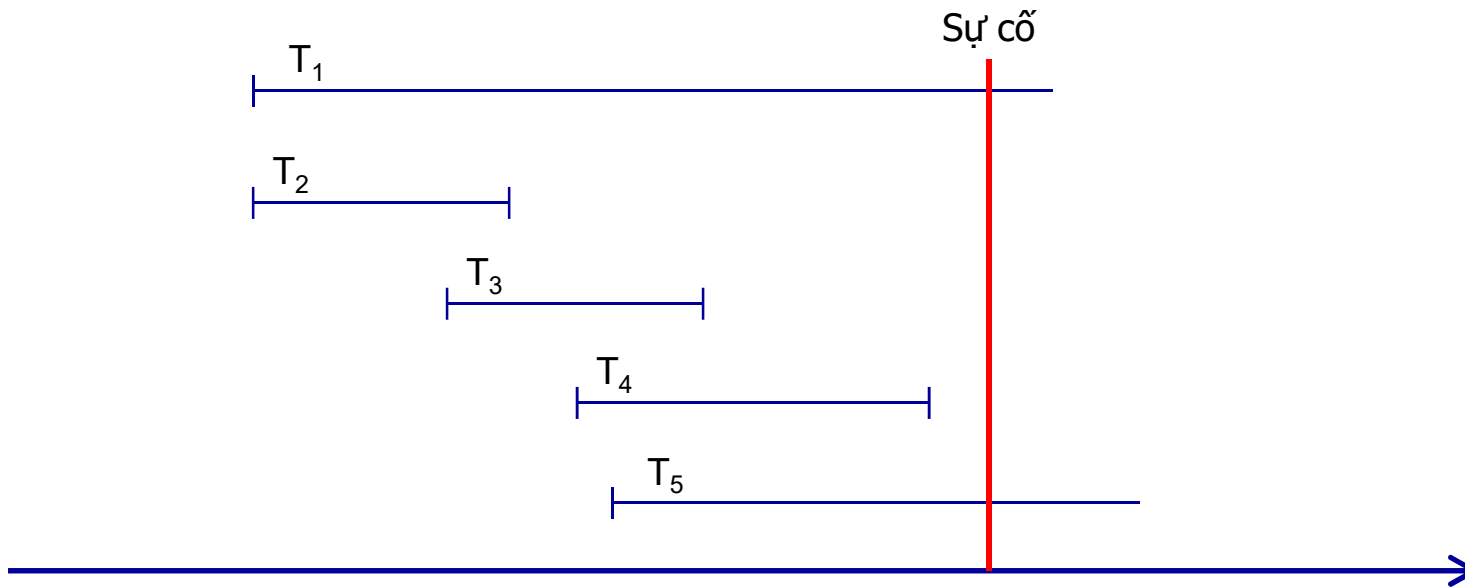


Note

- ☐ Data in buffer is flushed to disk in some cases, such as:
 - ☐ A particular command (eg. Commit).
 - ☐ When buffer get full.
- ☐ Implicit flush: force-writing.
- ☐ If failure occurs before data flush:
 - ☐ For transactions already committed, RM redo their actions (rollforward) for Durability.
 - ☐ For transactions not yet committed, RM undo their actions (rollback) for Atomicity.



Example



When failure occurs, T_2 , T_3 , T_4 has committed \rightarrow RM redo their data manipulation when DB restarts.

T_1 and T_5 will rollback





- DBMS provides these Recovery utilities:
 - RM.
 - Backup.
 - Logging.
 - Checkpoint.





Steal & No-force

- RM use 2 methods for flushing data from buffer to disk:
 - Steal policy: data in buffer is flushed to disk **Before transaction commit**. V.s no-steal, flushing nothing before transaction commit.
 - Force policy: data in buffer is flushed to disk immediately when **Transaction commit**. V.s no-force.
- For no-steal → No need to undo changes by uncommitted transactions.
- For force → No need to redo changes by committed transactions.
- Steal policy helps keep buffer from being full of need-flushing blocks.
- No force is beneficial when 2 transactions access the same block, the latter will not have to reload the block from disk.
- Most DBMS use: steal, no-force.





Log file

- Log file audits all changes made to DB.
- Log file is used for DB Recovery, consists of:
 - Transaction record:
 1. ID of transactions.
 2. Log records(Transaction start, insert, update, delete, abort, commit)
 3. ID of data items.
 4. Data items' old values \Leftrightarrow Before Image
 5. Data items' new values \Leftrightarrow After Image
 6. Pointer managing log records in log file.
 - Checkpoint record.
- Because of log files' extreme importance, DBMS often maintains 2 or 3 copy of log files simultaneously.



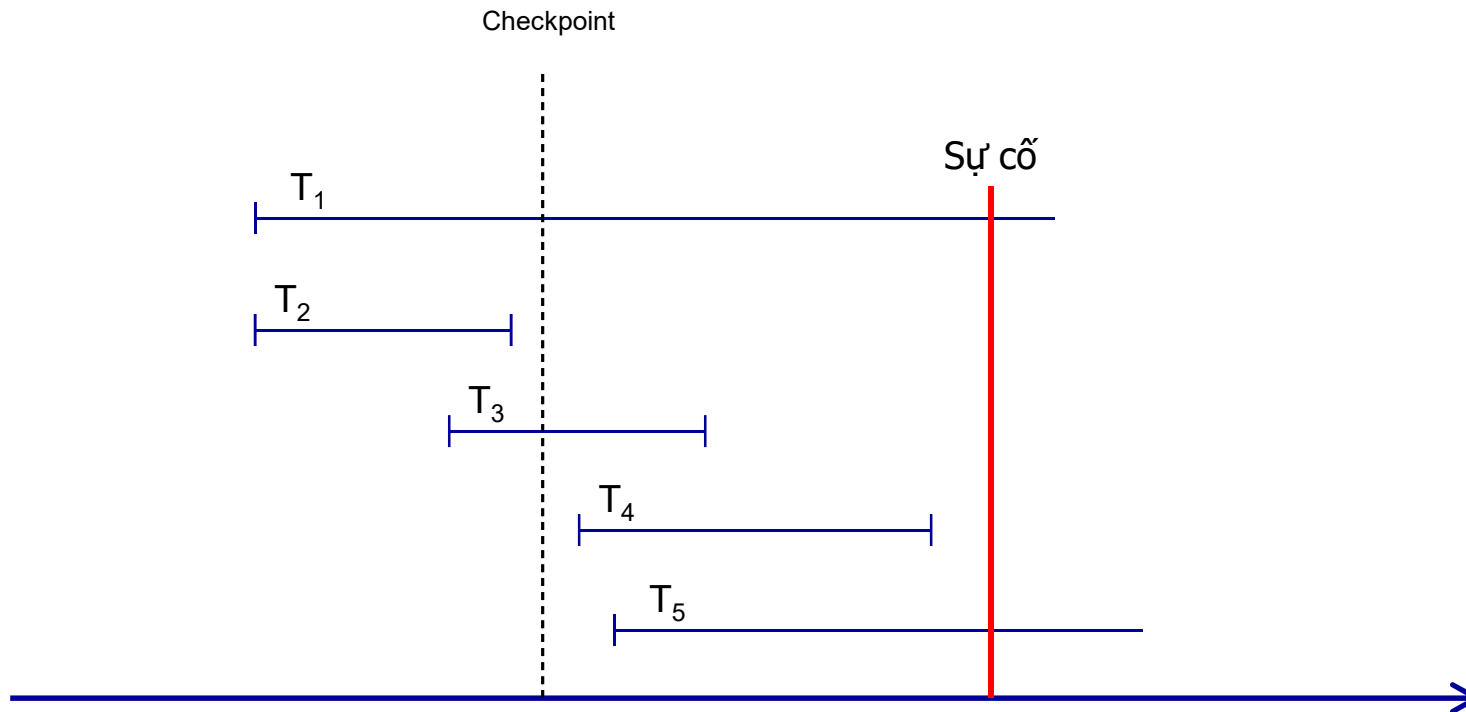


Checkpoint

- Recovery using log files has a drawback:
 - Scan the entire log → time consuming.
 - Unnecessary process of log records for already-written data blocks.
 - Checkpoint is introduced to address this disadvantage.
- Checkpoint will periodically:
 - Flush to log files all un-written log records in Log buffer.
 - Flush to disk all un-written data blocks in DB buffer.
 - Write checkpoint record to log file.
- RM decides the interval for checkpoint, may be after n minutes or when t transactions have committed after the previous checkpoint.



Checkpoint example



T_2 has been flushed, no need to redo T_2 .





Recovery techniques

- ☐ Recovery techniques using deferred update
- ☐ Recovery techniques using immediate update





Recovery techniques using deferred update

- ☐ No flushing until transaction commit. (no-steal.)
- ☐ Transaction failure before commit → No undo.
- ☐ Redo committed transactions.
- ☐ For transaction T, use log as:
 - ☐ Transaction T start, note to log.
 - ☐ No changes made to DB buffer or disk during transaction T.
 - ☐ Write all log records and commit record for T to log file.
 - ☐ Use log records to actually manipulate DB.
 - ☐ If T aborts, ignore all its log records, doing nothing to the DB.





Recovery techniques using immediate update

- ☐ Execute every transaction command immediately, does not wait until commit.
- ☐ When failure occurs:
 - ☐ Redo committed transactions.
 - ☐ Undo uncommitted transactions.
- ☐ For transaction T, use log file as:
 - ☐ T starts, note to log file.
 - ☐ Write every log records for T to log file.
 - ☐ Data changes are stored in DB buffer and flushed to disk when appropriate.
 - ☐ If T commit, note to log file.





WAL Protocol (Write Ahead Log)


- When failure, unflushed log records in log buffer may be lost in the same way as unflushed data blocks in DB buffer.
- Log records must be written to log file before corresponding data changes are flushed to DB on disk (Write Ahead Log Protocol).





UNDO protocol


- Undo uncommitted transactions or rolledback transactions





REDO Protocol

- Redo committed transactions which are not yet flushed to disk.





Normal recovery

- After a normal system shutdown, 1 checkpoint is written to the end of log file.
- When system starts, thanks to this checkpoint at the end of log file, no Undo or Redo is needed .





Failure recovery

- ☐ If the last record in log file is checkpoint, STOP.
- ☐ Else, find the last checkpoint in log file.
- ☐ Identify 2 transaction group:
 - ☐ Group 1: Committed transaction.
 - ☐ Group 2:
 - ☐ Uncommitted transactions.
 - ☐ Rollbacked transactions.
- ☐ For group 1: Redo.
- ☐ For group 2: Undo.





Convention

- Undo actions with ↑
 - ↑ Need actual undo on DB using before image.
 - [↑] No need actual undo on DB because actions take place after last checkpoint, changes are only in log file.
- Redo actions with ×



Ex:

BOT_i	Bắt đầu giao tác T_i
$U1(i)$	Cập nhật lần 1 của T_i
BOT_{i+1}	Bắt đầu giao tác T_{i+1}
$U1(i+1)$ ↑	Thao tác cập nhật thứ 1 của giao tác T_{i+1}
<i>Checkpoint</i>	
BOT_{i+2}	Bắt đầu giao tác T_{i+2}
$U1(i+2)$ ✕	Thao tác cập nhật thứ 1 của giao tác T_{i+2}
$U2(i)$ ✕	Thao tác cập nhật thứ 2 của giao tác T_i
Commit T_i	Commit T_i
$U2(i+1)$ ↗	Thao tác cập nhật thứ 2 của giao tác T_{i+1}
BOT_{i+3}	Bắt đầu giao tác T_{i+3}
$U1(i+3)$ ↗	Thao tác cập nhật thứ 1 của giao tác T_{i+3}
$U2(i+3)$ ↗	Thao tác cập nhật thứ 2 của giao tác T_{i+3}
$U2(i+2)$ ✕	Thao tác cập nhật thứ 2 của giao tác T_{i+2}
Commit T_{i+2}	Commit T_{i+2}
$U3(i+1)$ ↗	Thao tác cập nhật thứ 3 của giao tác T_{i+1}
Sự cố hệ thống xảy ra	


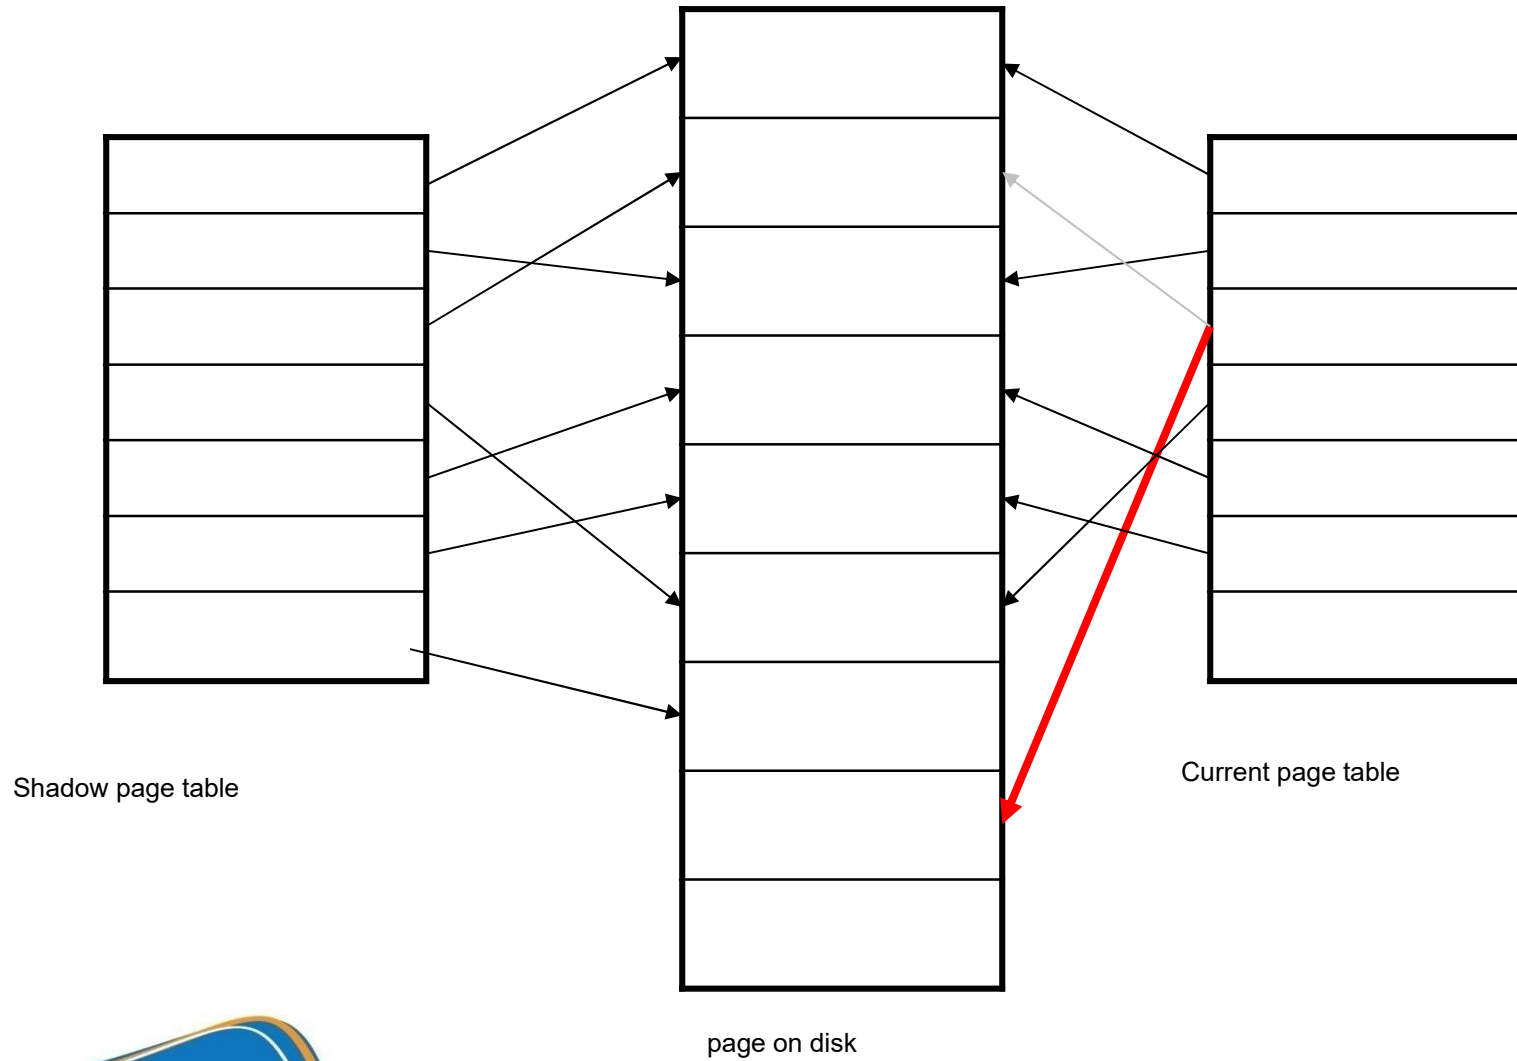


Shadow paging

- Another Recovery method is Shadow paging.
- During transaction T's entire life time, 2 tables are maintained:
 - ▣ Current page table: Changed as T writes data.
 - ▣ Shadow page table: A copy of the table before T starts.
 - ▣ When T starts, the 2 page tables are the same.



Shadow paging



End of chapter 3. Part 1

