

Chapter 3

Part 1 DATABASE RECOVERY TECHNIQUES

1

1

Outline

1. Goals
2. Type of Failures
3. Major Concepts
 - Transaction log
 - Checkpoint
4. Database Recovery Techniques

2

2

Example

- ☐ Accounts: A : 1000 \$, B : 2000 \$
- ☐ Transfer 50 \$ from A to B.
- ☐ Let's consider the state:
 - ☐ A had updated: $A := A - 50$
 - ☐ B had not been updated: $B := B + 50$
 - ☐ And there was power cut!
- ☐ When the system restarted:
 - ☐ If T was executed again: $A = 900$
 - ☐ If not: $A = 950$ and $B = 2000$.
 - ☐ How should the system be recovered?

3

3

Goals

- ☐ Database Recovery is the process of restoring the database and the data to a consistent state. This may include restoring lost data up to the point of the event (e.g. system crash).
- ☐ This is done by RM – the Recovery Manager.
- ☐ Automatic database recovery helps saving manual effort after every crash.

4

4

Goals

- ☐ Transaction is the basic unit of Database Recovery.
- ☐ Among ACID properties of Transaction, Database Recovery helps retain Atomic and Durability.

5

5

Types of failures

- ☐ Transaction failures
 - ☐ Rolling back caused by Deadlock or as requested by the scheduler.
 - ☐ During failure, system still operates normally.
 - ☐ Frequency: some times/minute.
- ☐ System failures
 - ☐ System can no longer operate. The reason may be failures in the processing unit, power cut or software failure.
 - ☐ Only cause data lost on RAM.
 - ☐ Frequency: some times/month.

6

6

Types of failures

- ☐ Media failure
 - ☐ Ex: HDD crash.
 - ☐ Cause the lost of a part or entire database.
 - ☐ Frequency: some times/year.
- ☐ Software failures
 - ☐ Logical error of applications which access database, leading to transaction failure.

7

7

Note

- ☐ Regardless of failure reasons, we have to consider 2 issues:
 1. Data lost on database buffer.
 2. Data lost on storage media.

8

8

Backup


- DBMS provide Backup mechanism.
 - ▣ Full backup / Differential backup.
 - ▣ Handle data lost on storage media.

9

9

DB access

- Data is read or written in primary unit called “block”.
 - ▣ Physical block: Stored in Media.
 - ▣ Buffer block: Stored in buffer.



10

10

DB access

- ☐ Read (X): assign X to local variable
 - ☐ If the data block with X is not in buffer, then Input (X).
 - ☐ assign X to local variable Xi.
- ☐ Write (X): assign Xi to X
 - ☐ If the data block with X is not in buffer, then Input (X)..
 - ☐ Gán giá trị xi cho X (trên buffer block có chứa X).

11

11

Buffer management

- ☐ Buffer:
 - ☐ Dữ liệu mất khi có sự cố hệ thống.
 - ☐ Không gian hạn chế.
- ☐ Chiến lược thay thế để định ra vùng trống trên buffer dùng để nạp dữ liệu mới.
 - ☐ FIFO.
 - ☐ LRU.


12

12

Note

- ☐ Database has 2 parts:
 - ☒ Physical Database and
 - ☒ Database Buffer

- ☐ Log has 2 parts:
 - ☒ Physical Log and
 - ☒ Log Buffer



- ☐ Failure \Rightarrow data lost in database buffer \Rightarrow use log file to recover data.
- ☐ Failure \Rightarrow data lost in log buffer \Rightarrow redo DB manipulation which has not been saved to physical DB.

13

13


Note

- ☐ Data in buffer is flushed to disk in some cases, such as:
 - ☒ A particular command (eg. Commit).
 - ☒ When buffer get full.
- ☐ Implicit flush: force-writing.
- ☐ If failure occurs before data flush:
 - ☒ For transactions already committed, RM redo their actions (rollforward) for Durability.
 - ☒ For transactions not yet committed, RM undo their actions (rollback) for Atomicity.

14

14

Example



When failure occurs, T₂, T₃, T₄ has committed → RM redo their data manipulation when DB restarts.

T₁ and T₅ will rollback

15

15

□ DBMS provides these Recovery utilities:

- ▣ RM.
- ▣ Backup.
- ▣ Logging.
- ▣ Checkpoint.

16

16

Steal & No-force

- ☐ RM use 2 methods for flushing data from buffer to disk:
 - ☐ Steal policy: data in buffer is flushed to disk **Befor transaction commit**. V.s no-steal, flushing nothing befor transaction commit.
 - ☐ Force policy: data in buffer is flushed to disk immediately when **Transaction commit**. V.s no-force.
- ☐ For no-steal → No need to undo changes by uncommitted transactions.
- ☐ For force → No need to redo changes by committed transactions.
- ☐ Steal policy helps keep buffer from being full of need-flushing blocks.
- ☐ No force is beneficial when 2 transactions access the same block, the latter will not have to reload the block from disk.
- ☐ Most DBMS use: steal, no-force.

17

17

Log file

- ☐ Log file audits all changes made to DB.
- ☐ Log file is used for DB Recovery, consists of:
 - ☐ Transaction record:
 1. ID of transactions.
 2. Log records(Transaction start, insert, update, delete, abort, commit)
 3. ID of data items.
 4. Data items' old values ⇔ Before Image
 5. Data items' new values ⇔ After Image
 6. Pointer managing log records in log file.
 - Checkpoint record.
- Because of log files' extreme importance, DBMS often maintains 2 or 3 copy of log files simultaneously.

18

18

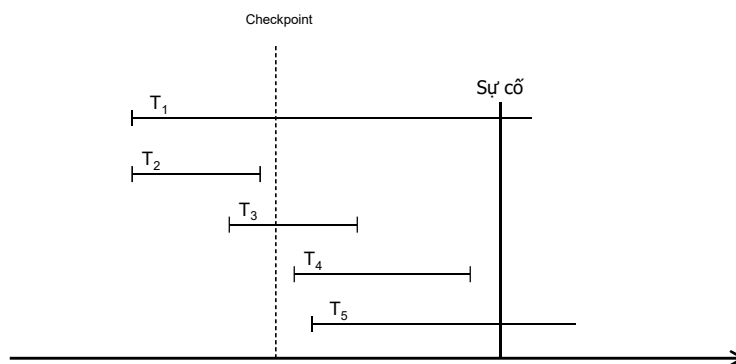
Checkpoint

- Recovery using log files has a drawback:
 - ▣ Scan the entire log → time consuming.
 - ▣ Unnecessary process of log records for already-written data blocks.
 - ▣ Checkpoint is introduced to address this disadvantage.
- Checkpoint will periodically:
 - ▣ Flush to log files all un-written log records in Log buffer.
 - ▣ Flush to disk all un-written data blocks in DB buffer.
 - ▣ Write checkpoint record to log file.
- RM decides the interval for checkpoint, may be after n minutes or when t transactions have committed after the previous checkpoint.

19

19

Checkpoint example



T2 has been flushed, no need to redo T2.

20

20

Recovery techniques

- ☐ Recovery techniques using deferred update
- ☐ Recovery techniques using immediate update

21

21

Recovery techniques using deferred update

- ☐ No flushing until transaction commit. (no-steal.)
- ☐ Transaction failure before commit → No undo.
- ☐ Redo committed transactions.
- ☐ For transaction T, use log as:
 - ▣ Transaction T start, note to log.
 - ▣ No changes made to DB buffer or disk during transaction T.
 - ▣ Write all log records and commit record for T to log file.
 - ▣ Use log records to actually manipulate DB.
 - ▣ If T aborts, ignore all its log records, doing nothing to the DB.

22

22

Recovery techniques using immediate update

- ☐ Execute every transaction command immediately, does not wait until commit.
- ☐ When failure occurs:
 - ☐ Redo committed transactions.
 - ☐ Undo uncommitted transactions.
- ☐ For transaction T, use log file as:
 - ☐ T starts, note to log file.
 - ☐ Write every log records for T to log file.
 - ☐ Data changes are stored in DB buffer and flushed to disk when appropriate.
 - ☐ If T commit, note to log file.

23

23

WAL Protocol (Write Ahead Log)


- ☐ When failure, unflushed log records in log buffer may be lost in the same way as unflushed data blocks in DB buffer.
- ☐ Log records must be written to log file before corresponding data changes are flushed to DB on disk (Write Ahead Log Protocol).

24

24

UNDO protocol

- Undo uncommitted transactions or rolledback transactions




25

25

REDO Protocol

- Redo committed transactions which are not yet flushed to disk.



26

26

Normal recovery

- ☐ After a normal system shutdown, 1 checkpoint is written to the end of log file.
- ☐ When system starts, thanks to this checkpoint at the end of log file, no Undo or Redo is needed .

27

27

Failure recovery

- ☐ If the last record in log file is checkpoint, STOP.
- ☐ Else, find the last checkpoint in log file.
- ☐ Indentify 2 transaction group:
 - ▣ Group 1: Committed transaction.
 - ▣ Group 2:
 - o Uncommitted transactions.
 - o Rollbacked transactions.
- ☐ For group 1: Redo.
- ☐ For group 2: Undo.

28

28

Convention

□ Undo actions with \uparrow

- ▣ \uparrow Need actual undo on DB using before image.
- ▣ $[\uparrow]$ No need actual undo on DB because actions take place after last checkpoint, changes are only in log file.

□ Redo actions with \times

29

29

Ex:

BOT_i	Bắt đầu giao tác T_i
U1(i)	Cập nhật lần 1 của T_i
BOT_{i+1}	Bắt đầu giao tác T_{i+1}
U1(i+1) \uparrow	Thao tác cập nhật thứ 1 của giao tác T_{i+1}
<i>Checkpoint</i>	
BOT_{i+2}	Bắt đầu giao tác T_{i+2}
U1(i+2) \times	Thao tác cập nhật thứ 1 của giao tác T_{i+2}
U2(i) \times	Thao tác cập nhật thứ 2 của giao tác T_i
Commit T_i	Commit T_i
U2(i+1) \uparrow	Thao tác cập nhật thứ 2 của giao tác T_{i+1}
BOT_{i+3}	Bắt đầu giao tác T_{i+3}
U1(i+3) \uparrow	Thao tác cập nhật thứ 1 của giao tác T_{i+3}
U2(i+3) \uparrow	Thao tác cập nhật thứ 2 của giao tác T_{i+3}
U2(i+2) \times	Thao tác cập nhật thứ 2 của giao tác T_{i+2}
Commit T_{i+2}	Commit T_{i+2}
U3(i+1) \uparrow	Thao tác cập nhật thứ 3 của giao tác T_{i+1}
	Sự cố hệ thống xảy ra

30

30

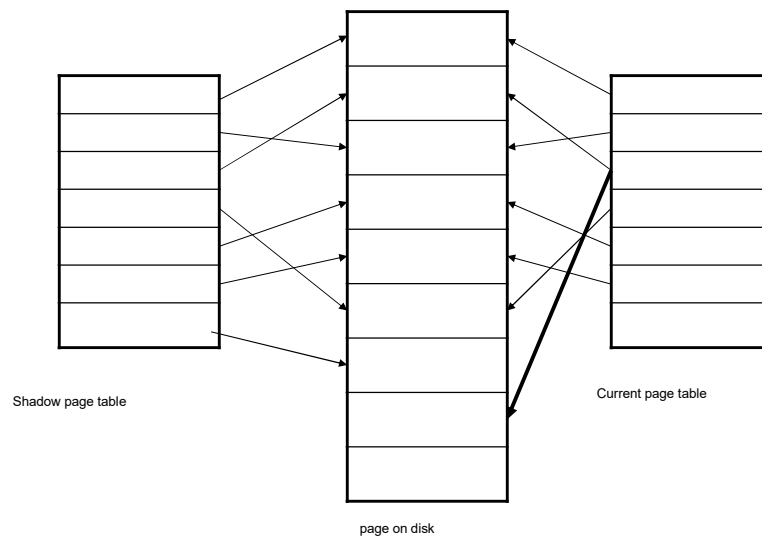
Shadow paging

- ☐ Another Recovery method is Shadow paging.
- ☐ During transaction T's entire life time, 2 tables are maintained:
 - ☒ Current page table: Changed as T writes data.
 - ☒ Shadow page table: A copy of the table before T starts.
 - ☒ When T starts, the 2 page tables are the same.

31

31

Shadow paging



32

32

End of chapter 3. Part 1

33

33